

Experiment 2

6 Docker Commands Are:

1. Build A Docker Image: Builds a Docker image from a Dockerfile

docker build -t [image_name] .

2. Run A Docker Container: Runs a container from an image

docker run -d -p 8080:80 --name [container_name] [image_name]

3. List All Containers: Lists all containers (running & stopped)

docker ps -a

4. Stop A Container: Stops a running container

docker stop [container_name]

5. Remove A Container: Deletes a stopped container

docker rm [container_name]

6. Remove A Image: Deletes a Docker image

docker rmi [image_name]

Experiment 3

Steps:

1. Create a empty folder.
2. Open VS Code in that folder.
3. Create Your index.html file (or whatever content you want) in your (working directory) folder.
4. Create a *Dockerfile*. If you want to run HTML files in your container then your dockerfile will be as follows:



```
Dockerfile > ...
1  FROM nginx:alpine
2
3  RUN rm -rf /usr/share/nginx/html/*
4
5  COPY index.html /usr/share/nginx/html/
6
7  EXPOSE 80
```

```
FROM nginx:alpine
RUN rm -rf /usr/share/nginx/html/*
COPY index.html /usr/share/nginx/html/
EXPOSE 80
```

5. Open VS code's integrated terminal.
6. Run the command "docker -v" to ensure that docker is up and running.
7. To build a image, run the following command:
docker build -t [image_name] .
8. Your image has been created. To confirm run
docker images
You should see your image name.
9. To create a container and run it, run this command:
docker run -d -p 8080:80 --name [container_name] [image_name]
10. To check if your container is running you can execute the command "docker ps -a" or you can go to localhost:8080 to check if that port is running the output of the container or not.
11. To stop the container you need to execute:
docker stop [container_name]
12. To remove the container you need to execute:
docker rm [container_name]

Experiment 4

Steps:

1. Open your terminal in your working directory.
2. Then, to pull nginx, run the following command:
docker pull nginx
3. To ensure that the image is pulled, you can run:
docker images
4. To create and run the container, execute the following command:
docker run -d -p 8080:80 --name [container_name] nginx
5. To check if your container is running, go to localhost:8080. You should see this page:

Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to nginx.org.
Commercial support is available at nginx.com.

Thank you for using nginx.

6. To stop a docker container, just run:
docker stop [container_name]

Experiment 5

Steps:

1. Go to KillerCoda And log in.
2. Open Kubernetes 1.32
3. The master and the worker node will already be created. To view these nodes, run:

kubectl get nodes

4. The 6 Commands in Kubernetes are:

- To Get Nodes:

kubectl get nodes

- To Get pods:

kubectl get pods

- To Create a Pod:

kubectl run [pod_name] --image=nginx

- To Get Pod details:

kubectl describe pod [pod_name]

- To expose pod as a service:

kubectl expose pod [pod_name] --port=80 --type=NodePort

- To Get service:

kubectl get svc

Experiment 6

Steps:

1. Go to KillerCoda and open a Kubernetes playground.
2. To create a pod, you need to create a pod YAML file. So, run the following command:

nano mypod.yaml

3. Paste the following configuration in that file:

```
apiVersion: v1
kind: Pod
metadata:
  name: mypod
spec:
  containers:
  - name: mycontainer
    image: nginx
    ports:
    - containerPort: 80
```

```
apiVersion: v1
kind: Pod
metadata:
  name: mypod
spec:
  containers:
  - name: mycontainer
    image: nginx
    ports:
    - containerPort: 80
```

4. To save and exit nano: Ctrl + X → Y → Enter
5. Deploy the pod:

kubectl apply -f mypod.yaml

6. To ensure pod is running, execute:

kubectl get pods

7. To get the pod IP Address, run:

kubectl get pod mypod -o wide

8. To get the logs:

kubectl logs mypod

9. To describe a pod:

kubectl describe pod mypod

Feature	Host Machine	Guest Machine
Control	Controls the entire hardware and manages resources for all VMs.	Operates within the limits set by the host, under the hypervisor's control.
Resource Allocation	Has full access to CPU, RAM, disk, and other resources.	Only has access to a <i>portion</i> of CPU, RAM, disk, and devices allocated by the host.
Performance	Native performance (direct access to hardware).	Slightly lower performance due to virtualization overhead.
Operating System (OS)	Runs the main OS installed directly on the hardware.	Can run any OS (Windows, Linux, etc.), independent of the host OS.
Hardware Access	Direct access to real hardware like GPU, printer, USB ports, etc.	Accesses virtualized or emulated versions of hardware (unless specific passthrough is configured).